


---

# PARTITION ON THE FLY

NINA BELYAVSKAYA



---

## WHO I AM

My name is Nina Belyavskaya

I worked with PostgreSQL  
as a fullstack developer  
and as DBA.



---

## PostgreSQL PARTITIONING

Table partitioning is a great feature in PostgreSQL.

- Speed up query execution
- Prevent table bloating
- Easy archive old data

---

## DECLARATIVE PARTITIONING

Since version 10 we have declarative partitioning,  
which is a good way to create and arrange partitioned tables.


```
CREATE TABLE measurement (  
    city_id int not null,  
    logdate date not null,  
    peaktemp int,  
    unitsales int )  
PARTITION BY RANGE (logdate);
```

```
CREATE TABLE measurement_y2024m02 PARTITION OF measurement  
FOR VALUES FROM ('2024-02-01') TO ('2024-03-01');
```

```
CREATE TABLE measurement_y2024m03 PARTITION OF measurement  
FOR VALUES FROM ('2024-03-01') TO ('2024-04-01');
```

---

**BUT... IF HUGE TABLE ALREADY EXISTS?**



## EXAMPLE

```
CREATE TABLE events (  
    event_id serial PRIMARY KEY,  
    event_place_id integer NOT NULL,  
    created_at timestamp,  
    comment text,  
    username name NOT NULL  
);
```

```
CREATE TABLE places (  
    place_id integer PRIMARY KEY,  
    place_name text  
);
```

```
ALTER TABLE events  
    ADD CONSTRAINT events_place_fk  
    FOREIGN KEY (event_place_id)  
    REFERENCES places (place_id);
```

```
CREATE INDEX events_place_idx  
    ON events(event_place_id);
```

```
CREATE FUNCTION events_username() RETURNS trigger  
LANGUAGE plpgsql  
AS $$
```

```
    BEGIN  
        NEW.username := current_user;  
        RETURN NEW;  
    END;
```

```
$$;
```

```
CREATE TRIGGER events_username BEFORE INSERT ON events  
    FOR EACH ROW EXECUTE FUNCTION events_username();
```

---


```
INSERT INTO places (place_id, place_name)
(SELECT id, format('Place %s', id)
FROM generate_series(1, 1000) AS id);
```

```
INSERT INTO events (event_place_id)
( SELECT floor(random()*1000+1)
FROM generate_series(1,1000000) );
```

```
INSERT INTO events (event_place_id, created_at)
( SELECT floor(random()*1000+1), created_at
FROM generate_series('2023-01-01',now(),'1 min') AS created_at );
```

## IMPORTANT DECISIONS

- What data do we need continuous, non-downtime access to?
- What data do we need to store, e.g. for analytical purposes?





---

## ASSUME

We want to:

- access to this year data  
(`created_at`  $\geq$  '2024-01-01')
- keep all data (including nulls)
- There is no FK that references `events`



events

---

## PREPARATION – PARTITIONED TABLE

```
CREATE TABLE events_new (  
    LIKE events INCLUDING DEFAULTS INCLUDING CONSTRAINTS )  
    PARTITION BY RANGE (created_at);
```

```
ALTER TABLE events_new ADD PRIMARY KEY(event_id, created_at);  
CREATE UNIQUE INDEX ON events_new(event_id, created_at);
```

---

## PREPARATION – PARTITIONED TABLE

```
CREATE INDEX events_new_place_idx ON events_new(event_place_id);
```

```
ALTER TABLE events_new  
  ADD CONSTRAINT events_new_place_fk FOREIGN KEY (event_place_id)  
  REFERENCES places (place_id);
```

```
CREATE TRIGGER events_username BEFORE INSERT ON events_new  
  FOR EACH ROW EXECUTE FUNCTION events_username();
```

```
ALTER TABLE events_new ALTER COLUMN created_at SET DEFAULT now();
```

## PREPARATION – PARTITIONED TABLE

```
\d events_new
```

Partitioned table "public.events\_new"

Column	Type	Collation	Nullable	Default
event_id	integer		not null	nextval('events_event_id_seq'::regclass)
event_place_id	integer		not null	
created_at	timestamp without time zone			now()
comment	text			
username	name		not null	

Partition key: RANGE (created\_at)

Indexes:

"events\_new\_event\_id\_created\_at\_idx" UNIQUE, btree (event\_id, created\_at)

"events\_new\_place\_idx" btree (event\_place\_id)

Foreign-key constraints:

"events\_new\_place\_fk" FOREIGN KEY (event\_place\_id) REFERENCES places(place\_id)

Triggers:

events\_username BEFORE INSERT ON events\_new FOR EACH ROW EXECUTE FUNCTION events\_username()

Number of partitions: 0

---

## PREPARATION – PARTITIONS FOR NULLS AND OLD VALUES

```
CREATE TABLE events_null ( LIKE events );  
ALTER TABLE events_null  
    ADD CONSTRAINT events_created_at_null CHECK ( created_at IS NULL );
```

```
CREATE TABLE events_old ( LIKE events );  
ALTER TABLE events_old  
    ADD CONSTRAINT events_created_at_old CHECK ( created_at < '2024-01-01' );
```

---

## PREPARATION – PARTITION FOR ACTUAL VALUES

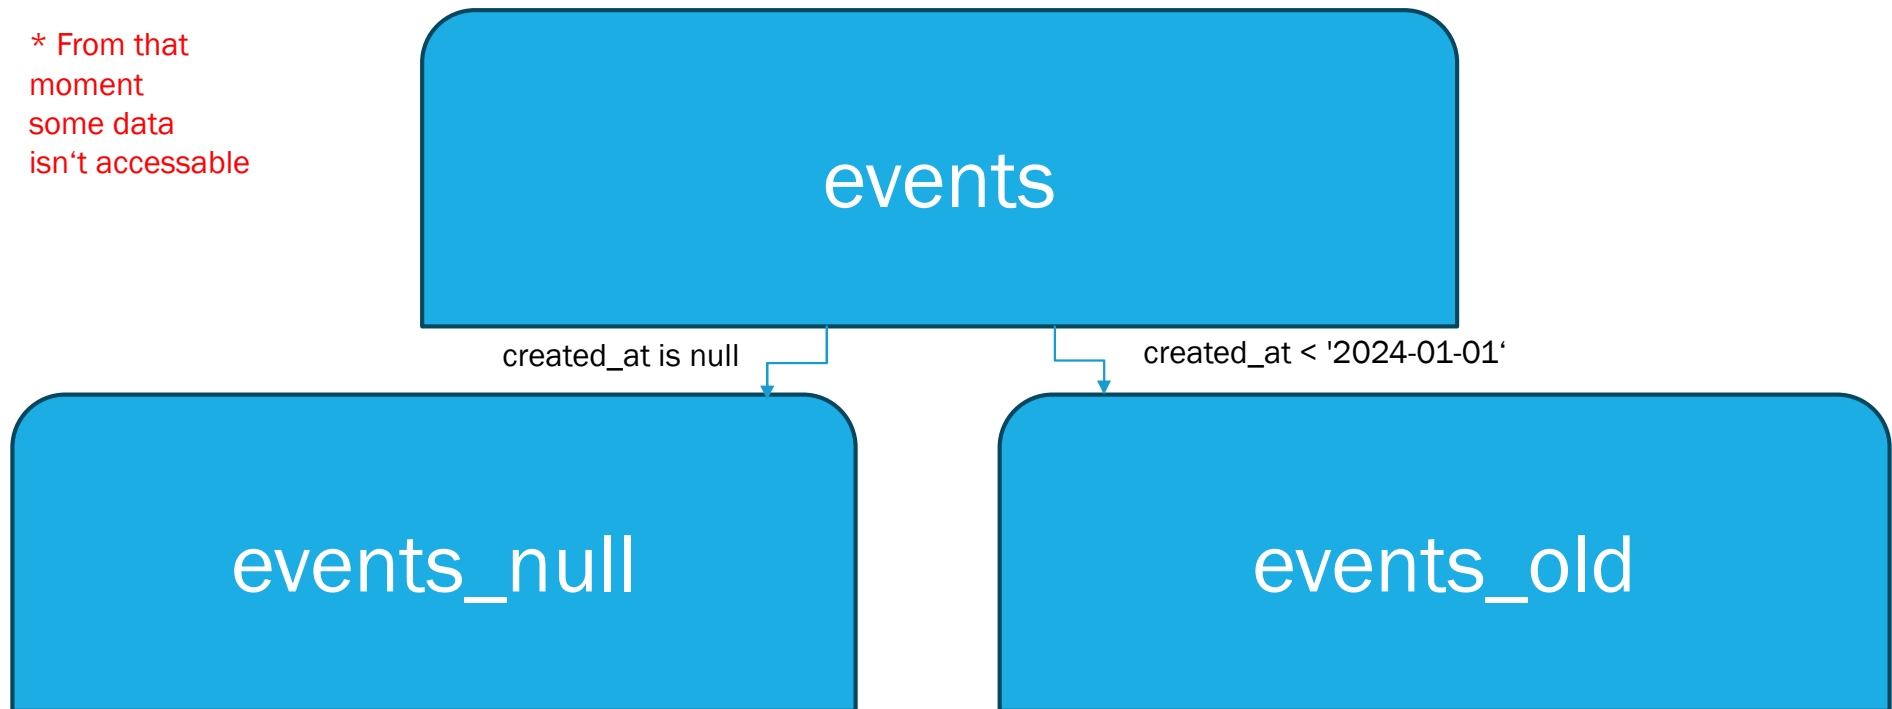
```
ALTER TABLE events ALTER COLUMN created_at SET DEFAULT now();
```

```
ALTER TABLE events  
  ADD CONSTRAINT events_created_at_not_null  
  CHECK ( created_at IS NOT NULL ) NOT VALID;
```

```
ALTER TABLE events  
  ADD CONSTRAINT events_created_at_2024  
  CHECK ( created_at >= '2024-01-01' ) NOT VALID;
```

## PREPARATION - MOVE OLD DATA TO OTHER TABLES

\* From that  
moment  
some data  
isn't accessible



---

## ATTACH PARTITIONS

- ALTER TABLE events\_new  
ATTACH PARTITION events\_null DEFAULT;
- ALTER TABLE events\_new  
ATTACH PARTITION events\_old  
FOR VALUES FROM (minvalue) TO ('2024-01-01');



---

## VALIDATE CONSTRAINTS

- ALTER TABLE events  
VALIDATE CONSTRAINT events\_created\_at\_not\_null;
- ALTER TABLE events  
VALIDATE CONSTRAINT events\_created\_at\_2024;

---

## PARTITIONED TABLE ASSEMBLY

**BEGIN;**

ALTER TABLE events RENAME TO events\_2024\_03;

ALTER TABLE events\_new RENAME TO events;

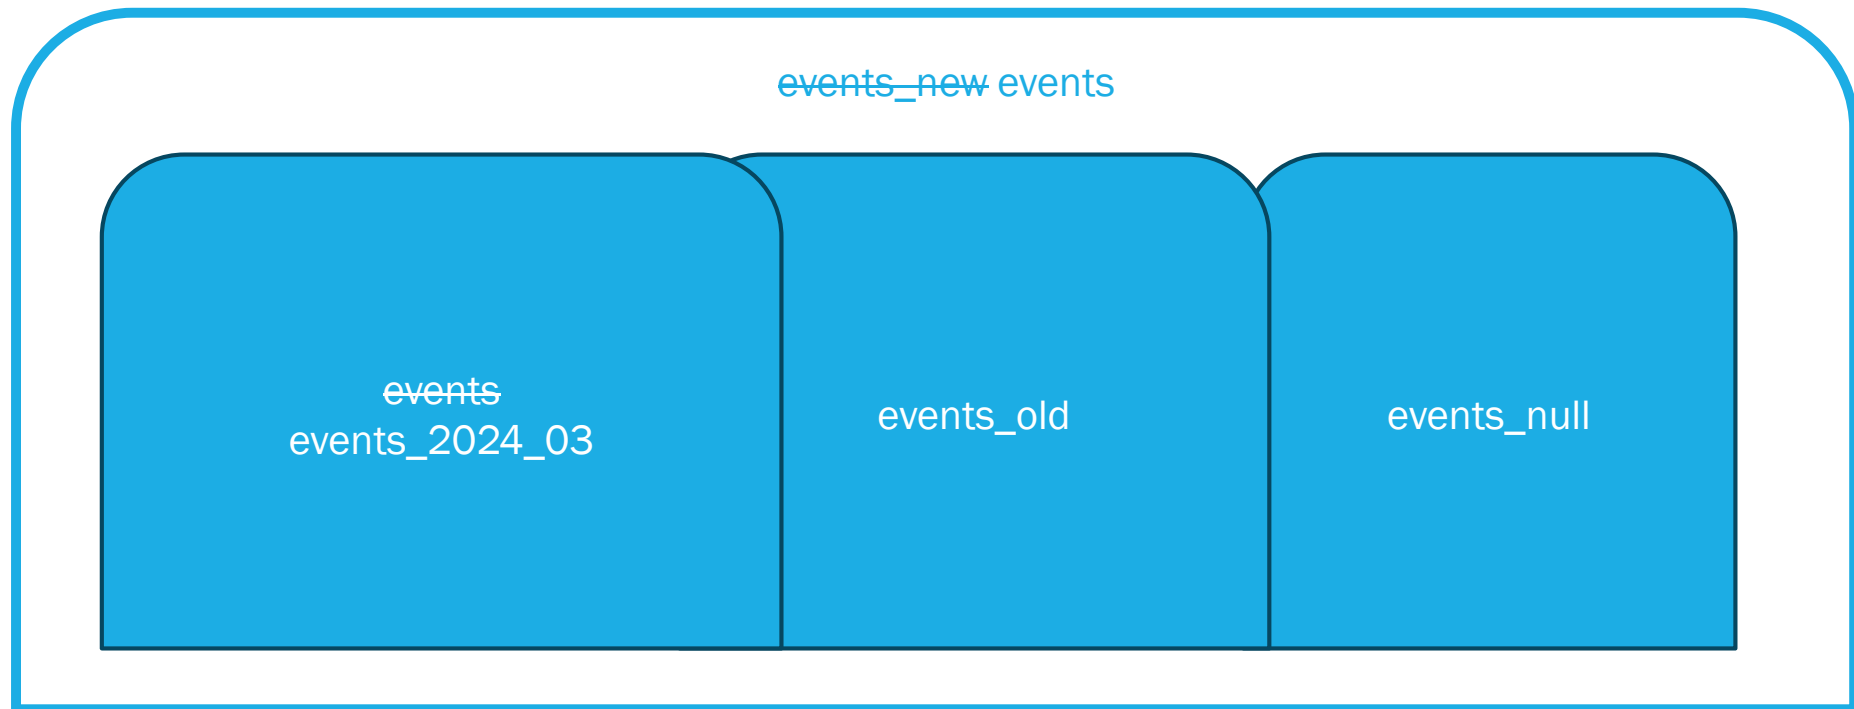
DROP TRIGGER events\_username ON events\_2024\_03;

ALTER TABLE events ATTACH PARTITION events\_2024\_03  
FOR VALUES FROM ('2024-01-01') TO ('2024-04-01');

**COMMIT;**

---

## PARTITIONED TABLE ASSEMBLY



Partitioned table "public.events"

Column	Type	Collation	Nullable	Default
event_id	integer		not null	nextval('events_event_id_seq'::regclass)
event_place_id	integer		not null	
created_at	timestamp without time zone			
comment	text			
username	name		not null	

Partition key: RANGE (created\_at)

Indexes:

"events\_new\_event\_id\_created\_at\_idx" UNIQUE, btree (event\_id, created\_at)

"events\_new\_place\_idx" btree (event\_place\_id)

Foreign-key constraints:

"events\_new\_place\_fk" FOREIGN KEY (event\_place\_id) REFERENCES places(place\_id)

Triggers:

events\_username BEFORE INSERT ON events FOR EACH ROW EXECUTE FUNCTION events\_username()

Partitions: events\_2024\_03 FOR VALUES FROM ('2024-01-01 00:00:00') TO ('2024-04-01 00:00:00'),

events\_old FOR VALUES FROM (MINVALUE) TO ('2024-01-01 00:00:00'),

events\_null DEFAULT



**Don't forget:**

```
\d events_event_id_seq
```

```
...
```

```
Owned by: public.events_2024_03.event_id
```

```
ALTER SEQUENCE events_event_id_seq  
OWNED BY events.event_id;
```

---

# Q & A

MY CONTACTS

<https://www.linkedin.com/in/nina-belyavskaya/>

[ninaobel@gmail.com](mailto:ninaobel@gmail.com)